

Credit Card Fraud Detection

In this project you will predict fraudulent credit card transactions with the help of Machine learning models. Please import the following libraries to get started.

```
import pandas
import numpy
import datetime

import warnings
import missingno as msno

pandas.set_option('display.max_columns', 500)
warnings.simplefilter('ignore')

import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib.colors import ListedColormap

from sklearn import preprocessing
from sklearn.decomposition import PCA
from sklearn.decomposition import IncrementalPCA

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import precision_recall_curve
from sklearn.feature_selection import RFE
from sklearn import metrics

import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.linear_model import BayesianRidge
from sklearn.impute import KNNImputer

from collections import Counter
from imblearn.over_sampling import SMOTE
from imblearn.over_sampling import ADASYN, RandomOverSampler
```

Utility Functions

```
def format_date(x):
    if str(x) == 'nan':
```

```

        return x
    else:
        return datetime.datetime.strptime(x, "%m/%d/%Y")

def imputation(data, column, value):
    data.loc[data[column].isnull()==True,column] = value
    return data

def missing_value_percentage(data):
    missing_value_summary =
pandas.DataFrame(data.isnull().sum()*100/data.shape[0])
    missing_value_summary.columns = ['Invalid Data %']
    missing_value_summary =
missing_value_summary.loc[missing_value_summary['Invalid Data
%']>0].sort_values('Invalid Data %', ascending=False)
    return missing_value_summary

def create_outlier_df(data, var_list):

outlier_df=pandas.DataFrame(columns=['ColumnName','OutlierCount','Outl
ier%'])
    for var in var_list:
        try:
            Q1 = data[var].quantile(0.25)
            Q3 = data[var].quantile(0.75)
            IQR = Q3 - Q1
            outlier_count=((data[var] < (Q1 - (1.5 * IQR))) |
(data[var] > (Q3 + (1.5 * IQR)))).sum()
            new_row = {
                'ColumnName':var,
                'OutlierCount':outlier_count,
                'Outlier
%':round(outlier_count*100/data[var].shape[0],2)
            }
            outlier_df=outlier_df.append(new_row,ignore_index=True)
        except TypeError:
            print('Error with column '+var)
    return outlier_df

def cap_outlier(data, var_list):
    for col in var_list:
        Q1 = data[col].quantile(0.25)
        Q3 = data[col].quantile(0.75)
        IQR = Q3 - Q1
        data[col][data[col] <= (Q1 - 1.5 * IQR)] = (Q1 - 1.5 * IQR)
        data[col][data[col] >= (Q3 + 1.5 * IQR)] = (Q3 + 1.5 * IQR)
    return data

```

```

def vif_ranks(data, features, row_count):
    vif = pandas.DataFrame()
    vif['Features'] = data[features].columns
    vif['VIF'] = [variance_inflation_factor(X_train[features].values,
i) for i in range(X_train[features].shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return vif.head(row_count)

```

```

def create_correlation_df(data, var_list, threshold):
    listi=[]
    listj=[]
    data=data[var_list]
    resultDf=pandas.DataFrame(columns=['Feature 1', 'Feature
2', 'Correlation Value'])
    corrDf=data.corr()
    for i in corrDf.columns:
        for j in corrDf.columns:
            if i==j:
                break
            if (corrDf.loc[i,j] >=threshold) and (str(i)!=str(j)):

                new_row = {
                    'Feature 1':str(i), 'Feature 2':str(j),
                    'Correlation Value':round(corrDf.loc[i,j],2)
                }
                resultDf=resultDf.append(new_row,ignore_index=True)
    return resultDf

```

Performance Metrics Functions

```

def predict_summarize(predicted, actual, threshold, plot_roc_=False):
    y_pred_final = pandas.DataFrame({'Converted':actual,
'Converted_Probability':predicted})
    y_pred_final['CustID'] = range(len(predicted))
    y_pred_final['predicted'] =
y_pred_final['Converted_Probability'].map(lambda x: 1 if x > threshold
else 0)

```

```

# Confusion matrix
confusion = metrics.confusion_matrix(y_pred_final['Converted'],
y_pred_final['predicted'] )

```

```

TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives

```

```

sensitivity = TP / float(TP+FN)
specificity = TN / float(TN+FP)

```

```

false_positive_rate = FP/ float(TN+FP)
precision = TP / float(TP + FP)
recall = TP / float(TP + FN)

confusion = pandas.DataFrame(confusion)
confusion.columns = ['predicted_no','predicted_yes']
confusion['ind'] = ['actual_no','actual_yes']
confusion = confusion.set_index('ind')

print('Accuracy =
',metrics.accuracy_score(y_pred_final['Converted'],
y_pred_final.predicted))
print('Sensitivity = ',sensitivity)
print('Specificity = ',specificity)
print('False Positive Rate = ',false_positive_rate)
print('\nPrecision = ',precision)
print('Recall = ',recall)

if plot_roc_:
    plot_roc(y_pred_final['Converted'],
y_pred_final['Converted_Probability'])

return confusion

def plot_roc( actual, probs ):
    print('Plotting')
    fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
                                                drop_intermediate =
False )
    auc_score = metrics.roc_auc_score( actual, probs )
    plt.figure(figsize=(5, 5))
    plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic example')
    plt.legend(loc="lower right")
    plt.show()

    return None

def plot_precision_recall_curve(predicted, actual):
    y_pred_final = pandas.DataFrame({'Converted':actual,
'Converted_Probability':predicted})
    y_pred_final['CustID'] = range(len(actual))

    p, r, thresholds =

```

```

precision_recall_curve(y_pred_final['Converted'],
y_pred_final['Converted_Probability'])
    plt.figure(figsize=(10,5))
    plt.plot(thresholds, p[:-1], "g-",label='precision')
    plt.plot(thresholds, r[:-1], "r-",label='recall')
    plt.xlabel('Thresholds')
    plt.title('Precision Recall Curve')
    plt.legend()
    plt.show()

def plot_feature_importance(features, importances):
    feature_importance = pandas.DataFrame(importances, features)
    feature_importance['abs_value'] = abs(feature_importance[0])
    feature_importance =
feature_importance.sort_values('abs_value',ascending=False)
    plt.figure(figsize=(10,10))

sns.barplot(feature_importance[0],feature_importance.index,palette='h
sl')
    plt.title('Feature Importances')

```

Data Preparation

```

df = pandas.read_csv('.../data//creditcard.csv')
df.head()

```

	Time	V1	V2	V3	V4	V5	V6
V7 \							
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
0.239599							
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
0.078803							
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
0.791461							
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
0.237609							
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921
0.592941							
	V8	V9	V10	V11	V12	V13	
V14 \							
0	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390	-
0.311169							
1	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095	-
0.143772							
2	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293	-
0.165946							
3	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757	-
0.287924							
4	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852	-
1.119670							

	V15	V16	V17	V18	V19	V20
V21 \						
0	1.468177	-0.470401	0.207971	0.025791	0.403993	0.251412 -
	0.018307					
1	0.635558	0.463917	-0.114805	-0.183361	-0.145783	-0.069083 -
	0.225775					
2	2.345865	-2.890083	1.109969	-0.121359	-2.261857	0.524980
	0.247998					
3	-0.631418	-1.059647	-0.684093	1.965775	-1.232622	-0.208038 -
	0.108300					
4	0.175121	-0.451449	-0.237033	-0.038195	0.803487	0.408542 -
	0.009431					

	V22	V23	V24	V25	V26	V27
V28 \						
0	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558 -
	0.021053					
1	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983
	0.014724					
2	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353 -
	0.059752					
3	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723
	0.061458					
4	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422
	0.215153					

	Amount	Class
0	149.62	0
1	2.69	0
2	378.66	0
3	123.50	0
4	69.99	0

#observe the different feature type present in the data
df.dtypes

Time	float64
V1	float64
V2	float64
V3	float64
V4	float64
V5	float64
V6	float64
V7	float64
V8	float64
V9	float64
V10	float64
V11	float64
V12	float64

```

V13      float64
V14      float64
V15      float64
V16      float64
V17      float64
V18      float64
V19      float64
V20      float64
V21      float64
V22      float64
V23      float64
V24      float64
V25      float64
V26      float64
V27      float64
V28      float64
Amount    float64
Class      int64
dtype: object

```

```
df.describe()
```

	Time	V1	V2	V3
V4 \				
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.919560e-15	5.688174e-16	-8.769071e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00

	V5	V6	V7	V8
V9 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	-1.552563e-15	2.010663e-15	-1.694249e-15	-1.927028e-16
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01

```

1.343407e+01
25%   -6.915971e-01 -7.682956e-01 -5.540759e-01 -2.086297e-01 -
6.430976e-01
50%   -5.433583e-02 -2.741871e-01  4.010308e-02  2.235804e-02 -
5.142873e-02
75%    6.119264e-01  3.985649e-01  5.704361e-01  3.273459e-01
5.971390e-01
max    3.480167e+01  7.330163e+01  1.205895e+02  2.000721e+01
1.559499e+01

```

	V10	V11	V12	V13
V14 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	1.768627e-15	9.170318e-16	-1.810658e-15	1.693438e-15
std	1.088850e+00	1.020713e+00	9.992014e-01	9.952742e-01
min	-2.458826e+01	-4.797473e+00	-1.868371e+01	-5.791881e+00
25%	-5.354257e-01	-7.624942e-01	-4.055715e-01	-6.485393e-01
50%	-9.291738e-02	-3.275735e-02	1.400326e-01	-1.356806e-02
75%	4.539234e-01	7.395934e-01	6.182380e-01	6.625050e-01
max	2.374514e+01	1.201891e+01	7.848392e+00	7.126883e+00

	V15	V16	V17	V18
V19 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	3.482336e-15	1.392007e-15	-7.528491e-16	4.328772e-16
std	9.153160e-01	8.762529e-01	8.493371e-01	8.381762e-01
min	-4.498945e+00	-1.412985e+01	-2.516280e+01	-9.498746e+00
25%	-5.828843e-01	-4.680368e-01	-4.837483e-01	-4.988498e-01
50%	4.807155e-02	6.641332e-02	-6.567575e-02	-3.636312e-03
75%	6.488208e-01	5.232963e-01	3.996750e-01	5.008067e-01
max	8.877742e+00	1.731511e+01	9.253526e+00	5.041069e+00

	V20	V21	V22	V23
V24 \				

count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	5.085503e-16	1.537294e-16	7.959909e-16	5.367590e-16
std	7.709250e-01	7.345240e-01	7.257016e-01	6.244603e-01
min	-5.449772e+01	-3.483038e+01	-1.093314e+01	-4.480774e+01
25%	-2.117214e-01	-2.283949e-01	-5.423504e-01	-1.618463e-01
50%	-6.248109e-02	-2.945017e-02	6.781943e-03	-1.119293e-02
75%	1.330408e-01	1.863772e-01	5.285536e-01	1.476421e-01
max	3.942090e+01	2.720284e+01	1.050309e+01	2.252841e+01

	V25	V26	V27	V28
Amount \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	1.453003e-15	1.699104e-15	-3.660161e-16	-1.206049e-16
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

Here we will observe the distribution of our classes.

We also see that there are no missing values and hence no missing value treatment is required.


```

classes=df['Class'].value_counts()
normal_share=classes[0]/df['Class'].count()*100
fraud_share=classes[1]/df['Class'].count()*100

# Create a bar plot for the number and percentage of fraudulent vs
non-fraudulent transcations
plt.bar(['Non Fraud', 'Fraud'],
[normal_share/(normal_share+fraud_share)*100,
fraud_share/(normal_share+fraud_share)*100])

<BarContainer object of 2 artists>

```




Observation:

The percentage of fraud transactions is as low as .17 %


```

# Create a scatter plot to observe the distribution of classes with
time
plt.figure(figsize=(15,5))
plt.scatter(df['Time'], df['Class'])
plt.show()

```



```
# Create a scatter plot to observe the distribution of classes with
Amount
plt.figure(figsize=(15,5))
plt.scatter(df['Amount'], df['Class'])
plt.show()
```



We can clearly observe that all the fraudulent transactions are of relatively lower amounts.


```
# Drop unnecessary columns
df = df.drop('Time', axis=1)
```

Observation:

The time variable can be dropped as it does not have any relevance in the occurrence of fraud transactions

```
sns.heatmap(df.corr())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f57873e4690>
```

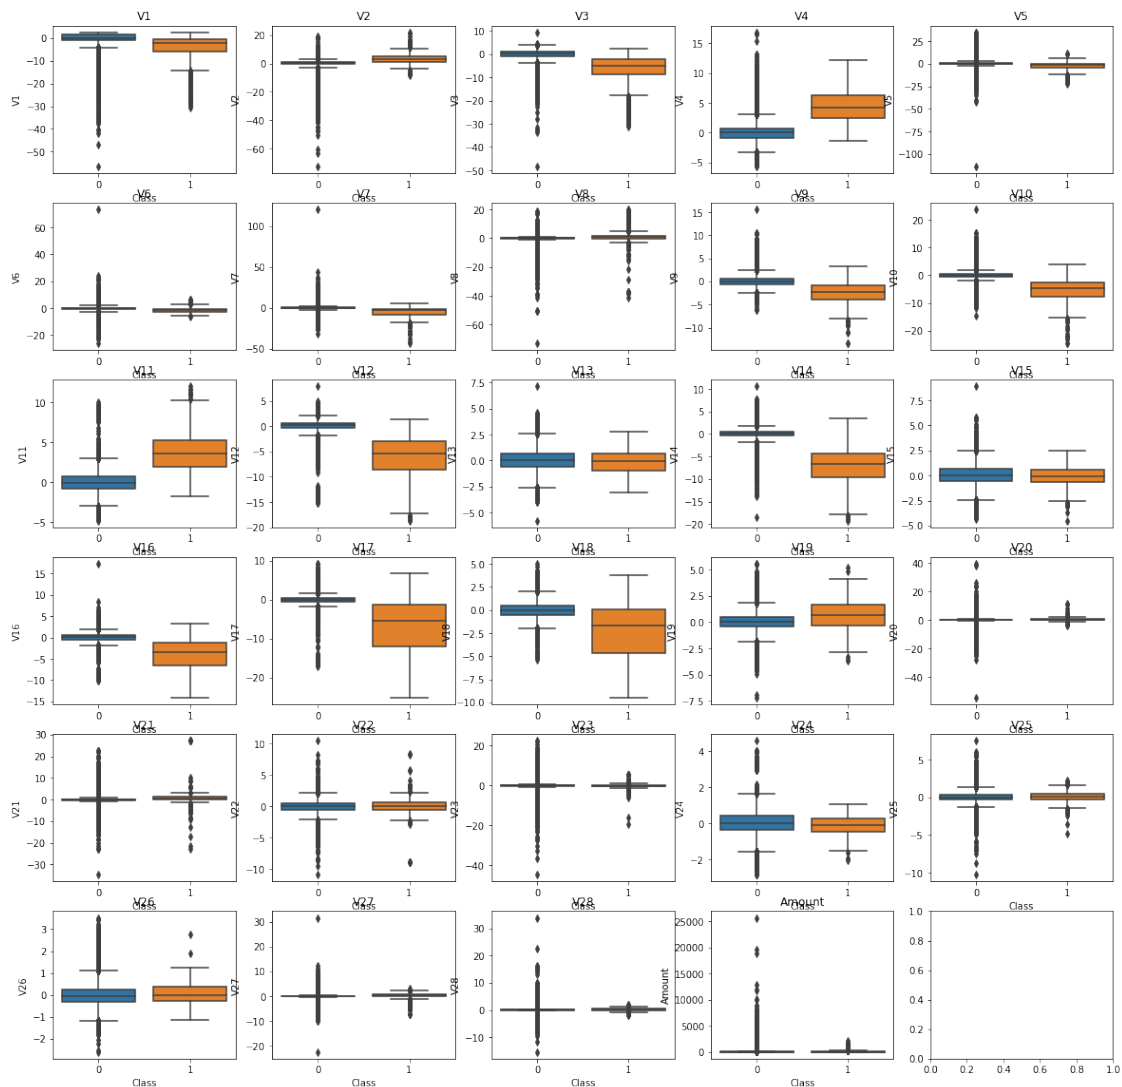


Observation:

No columns are correlated to one other, there we need not eliminate the features now.

```
column_list = ['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9',
               'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19',
               'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28',
               'Amount']
```

```
plt.subplots(6,5, figsize=(20,20))
index=1
for column in column_list:
    plt.subplot(6,5, index)
    sns.boxplot(y=df[column], x=df['Class'])
    plt.title(column)
    index += 1
plt.show()
```



Observation:

The columns do have a lot of outliers but in the case of mining out fraud transactions we can retain the outlier values as it.

Splitting the data into train & test data

Feature Analysis - Transformation

```
y = df.pop('Class') #class variable
X = df
```

```
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.3,
random_state=100)
```

```
print(numpy.sum(y))
print(numpy.sum(y_train))
print(numpy.sum(y_test))
```

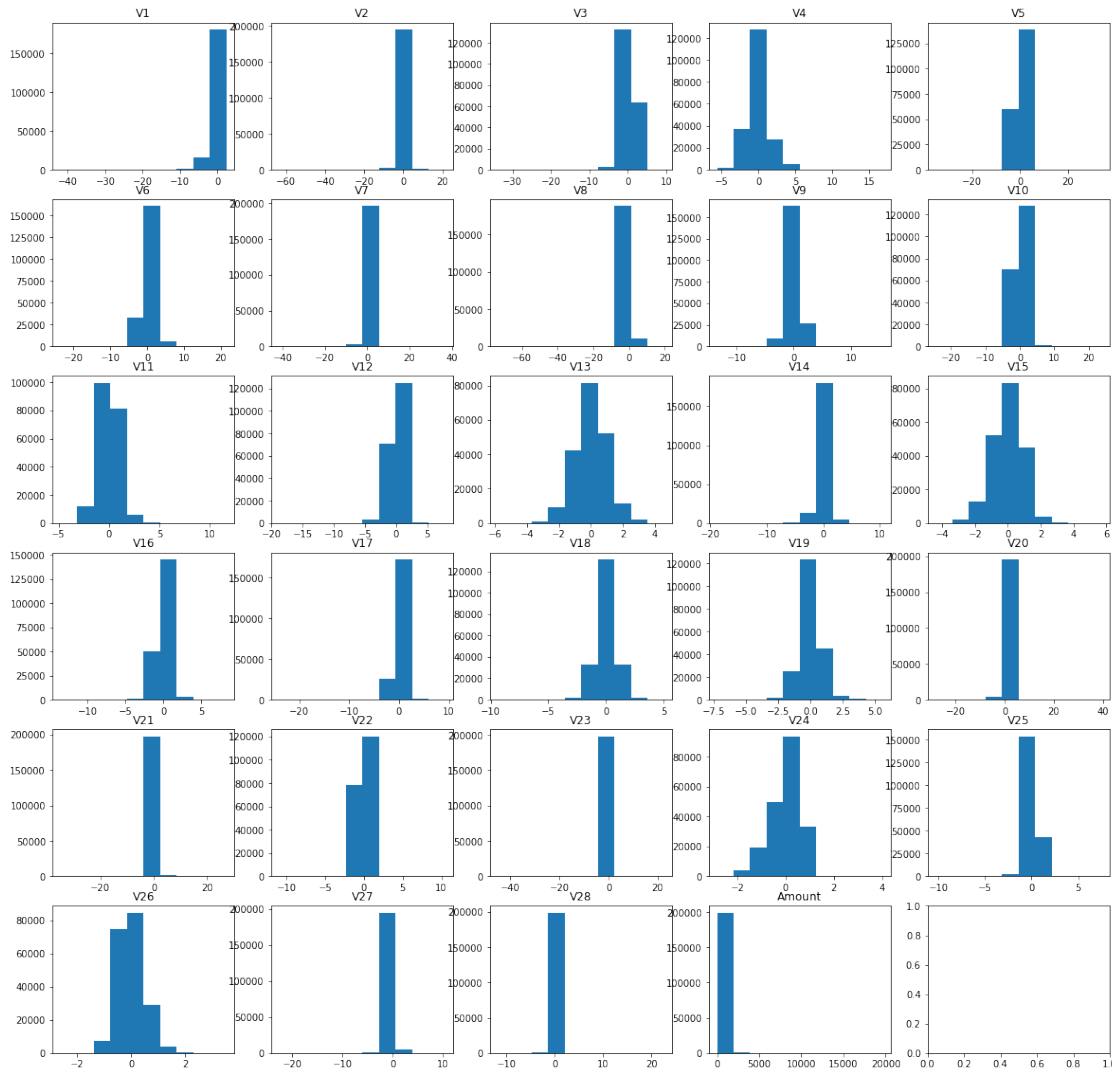
```
print('Train Fraud Rate', numpy.sum(y_train)/len(y_train))
print('Test Fraud Rate', numpy.sum(y_test)/len(y_test))

492
350
142
Train Fraud Rate 0.001755582753155033
Test Fraud Rate 0.001661926664559999
```

Observation:

The train and test sets both have similar event rates.

```
# plot the histogram of a variable from the dataset to see the
skewness
plt.subplots(6,5, figsize=(20,20))
index=1
for column in column_list:
    plt.subplot(6,5, index)
    plt.hist(X_train[column])
    plt.title(column)
    index += 1
plt.show()
```



Most of the above variables are skewed. So we can transform them to make them gaussian

- Apply : preprocessing.PowerTransformer(copy=False) to fit & transform the train & test data

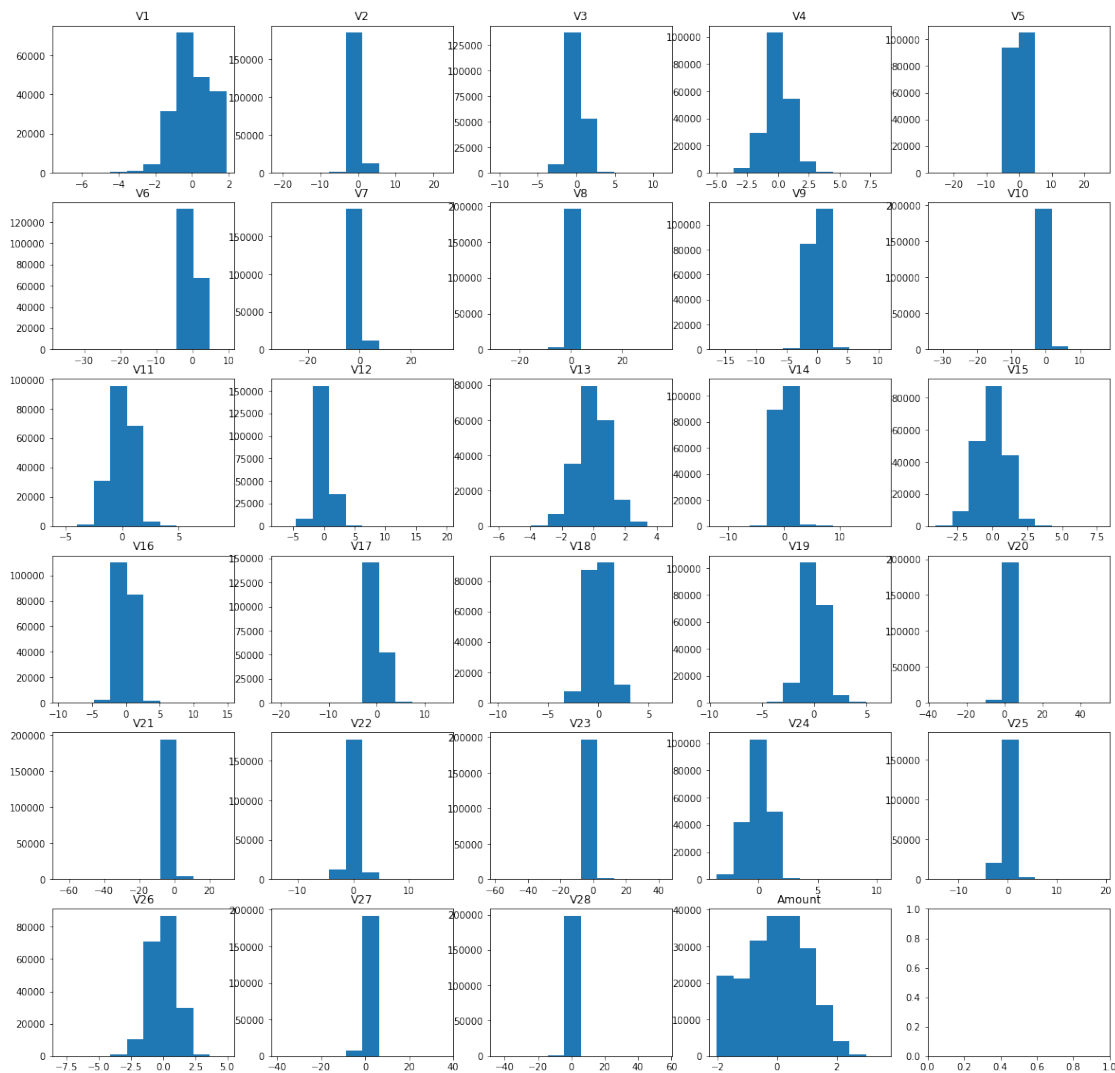
```
pt = preprocessing.PowerTransformer()
X_train[column_list] = pt.fit_transform(X_train[column_list])
X_test[column_list] = pt.transform(X_test[column_list])
```

```
scaler = preprocessing.StandardScaler()
X_train[X_train.columns] = scaler.fit_transform(X_train)
X_test[X_train.columns] = scaler.transform(X_test[X_train.columns])
```

plot the histogram of a variable from the dataset again to see the result

```
plt.subplots(6,5, figsize=(20,20))
index=1
for column in column_list:
    plt.subplot(6,5, index)
    plt.hist(X_train[column])
```

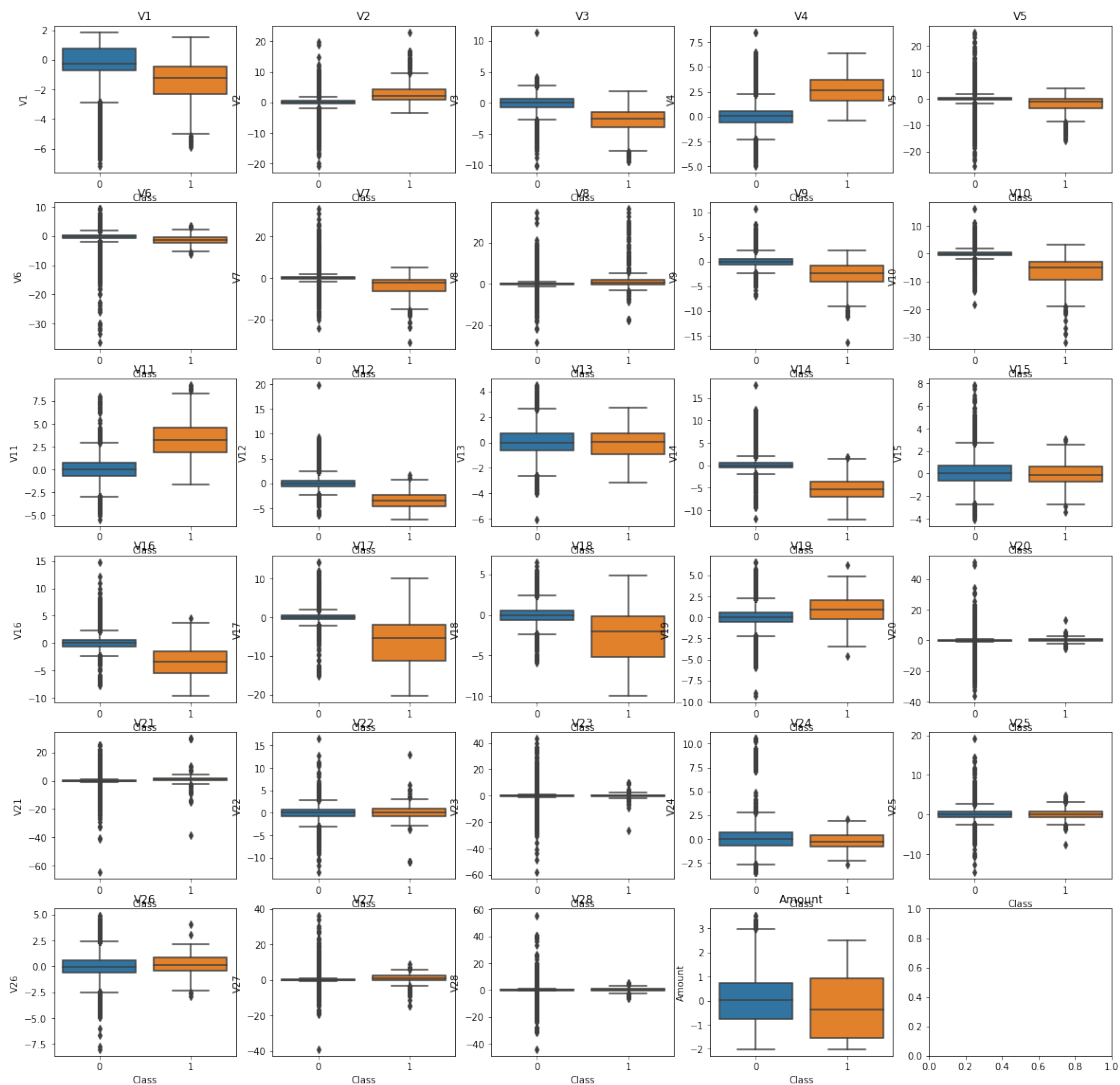
```
plt.title(column)
index += 1
plt.show()
```



Observation:

Few columns like amount column have undergone the transformation and have a different distribution now. The transformation was fit on train data and the test data was only transformed.

```
plt.subplots(6,5, figsize=(20,20))
index=1
for column in column_list:
    plt.subplot(6,5, index)
    sns.boxplot(y=X_train[column], x=y_train)
    plt.title(column)
    index += 1
plt.show()
```

Eliminate Insignificant Variables Using VIF

Though we have PCA transformed data, the objective still remains to be able to extract the important features impacting the churn. We will thus work towards this by eliminating the redundant variables.

```
features_set_1 = X_train.columns
vif_ranks(X_train, features_set_1, 10)
```

	Features	VIF
28	Amount	1.72
1	V2	1.60
0	V1	1.51
2	V3	1.28
4	V5	1.21
6	V7	1.10
7	V8	1.10
11	V12	1.09

```
5          V6    1.06
3          V4    1.04
```

The variables have $VIF < 3$. Therefore, we need not eliminate the variables at this stage.

Model Building

- Build different models on the imbalanced dataset and see the result

```
kf = KFold(n_splits=5, shuffle=False)
```

We can use kfold cross validation

Logistic Regression

```
'''
grid_params = {
    "C":numpy.logspace(-3, 3,7),
    "penalty":["l1", "l2"]
}
logreg = LogisticRegression()
logreg_cv = GridSearchCV(logreg, grid_params, cv=kf)
logreg_cv.fit(X_train, y_train)
print("tuned hpyerparameters :(best parameters) ",
logreg_cv.best_params_)
print("accuracy :", logreg_cv.best_score_)
'''


'\ngrid_params = {\n    "C":numpy.logspace(-3, 3,7), \n    "penalty":
["l1", "l2"]\n}\nlogreg = LogisticRegression()\nlogreg_cv =
GridSearchCV(logreg, grid_params, cv=kf)\nlogreg_cv.fit(X_train,
y_train)\nprint("tuned hpyerparameters :(best parameters) ",
logreg_cv.best_params_)\nprint("accuracy :", logreg_cv.best_score_)\n'

logreg = LogisticRegression(C=1, penalty='l2')

logreg.fit(X_train, y_train)

LogisticRegression(C=1, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

plot_precision_recall_curve(logreg.predict(X_train), y_train)
```



```
predict_summarize(logreg.predict(X_train), y_train, 0.8, True)
```

Accuracy = 0.9993529423566943

Sensitivity = 0.7057142857142857


Specificity = 0.9998693559247088

False Positive Rate = 0.00013064407529118555

Precision = 0.9047619047619048

Recall = 0.7057142857142857

Plotting



	predicted_no	predicted_yes
ind		
actual_no	198988	26
actual_yes	103	247

Observation:

This has a low recall value. We can continue with other models & recursive feature eliminations.

Logistic Regression - Iterations

```
rfe_feature_set_1=X_train.columns
```

```
features=rfe_feature_set_1
```

```
X_train_sm = sm.add_constant(X_train[features])
```

```
X_test_sm = sm.add_constant(X_test[features])
```

```
logm2 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
```

```
model = logm2.fit()
```

```
model.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

Generalized Linear Model Regression Results

```

=====
=====
Dep. Variable:          Class    No. Observations:
199364
Model:                  GLM      Df Residuals:
199334
Model Family:          Binomial  Df Model:
29
Link Function:         logit     Scale:
1.0000
Method:                IRLS      Log-Likelihood:
-672.87
Date:                  Sun, 20 Sep 2020    Deviance:
1345.7
Time:                  20:21:48    Pearson chi2:
4.99e+05
No. Iterations:        12

Covariance Type:        nonrobust

```

```

=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

const	-8.9200	0.183	-48.667	0.000	-9.279
-8.561					
V1	-0.0667	0.098	-0.683	0.494	-0.258
0.125					
V2	-0.2710	0.083	-3.248	0.001	-0.435
-0.107					
V3	-0.4474	0.099	-4.499	0.000	-0.642
-0.253					
V4	0.8344	0.096	8.660	0.000	0.646
1.023					
V5	-0.0743	0.081	-0.912	0.362	-0.234
0.085					
V6	0.1365	0.094	1.447	0.148	-0.048
0.321					
V7	-0.1907	0.070	-2.719	0.007	-0.328
-0.053					
V8	-0.2512	0.059	-4.279	0.000	-0.366
-0.136					
V9	-0.2690	0.099	-2.723	0.006	-0.463
-0.075					
V10	-0.1099	0.094	-1.174	0.241	-0.294
0.074					

V11	0.0054	0.097	0.056	0.955	-0.185
0.196					
V12	-0.7616	0.121	-6.318	0.000	-0.998
-0.525					
V13	-0.2510	0.096	-2.617	0.009	-0.439
-0.063					
V14	-0.8387	0.083	-10.045	0.000	-1.002
-0.675					
V15	-0.2010	0.096	-2.093	0.036	-0.389
-0.013					
V16	-0.3660	0.096	-3.827	0.000	-0.553
-0.179					
V17	0.0084	0.061	0.138	0.890	-0.111
0.128					
V18	0.1215	0.101	1.200	0.230	-0.077
0.320					
V19	-0.1316	0.087	-1.520	0.128	-0.301
0.038					
V20	-0.0761	0.057	-1.327	0.184	-0.188
0.036					
V21	0.2008	0.070	2.879	0.004	0.064
0.338					
V22	0.3241	0.112	2.889	0.004	0.104
0.544					
V23	-0.0717	0.043	-1.658	0.097	-0.157
0.013					
V24	0.0445	0.105	0.422	0.673	-0.162
0.251					
V25	0.1090	0.084	1.297	0.195	-0.056
0.274					
V26	-0.0680	0.117	-0.582	0.561	-0.297
0.161					
V27	-0.0711	0.055	-1.301	0.193	-0.178
0.036					
V28	-0.0580	0.033	-1.759	0.079	-0.123
0.007					
Amount	-0.0197	0.106	-0.187	0.852	-0.227
0.188					

```
=====
=====
"""
```

Observation:

We can iteratively eliminate the features with $p > 0.05$. The cell below was not executed at once, the set of features to be eliminated was done repeated basis until all vairables had $p < 0.05$

```
# removing the features with p-value >0.05, iteratively, one at a time
features=list(set(rfe_feature_set_1)-
set(['V11', 'V17', 'Amount', 'V24', 'V26', 'V1', 'V5', 'V18',
```

```
'V23','V25','V20','V28','V10','V27','V6','V19']))
```

```
X_train_sm = sm.add_constant(X_train[features])
X_test_sm = sm.add_constant(X_test[features])
logm2 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
model = logm2.fit()
model.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Class    No. Observations:
199364
Model:                  GLM      Df Residuals:
199350
Model Family:          Binomial  Df Model:
13
Link Function:          logit    Scale:
1.0000
Method:                 IRLS     Log-Likelihood:
-680.66
Date:                   Sun, 20 Sep 2020    Deviance:
1361.3
Time:                   20:21:49    Pearson chi2:
5.79e+05
No. Iterations:         12

Covariance Type:        nonrobust
```

```
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
const        -8.8567      0.162    -54.691      0.000     -9.174
-8.539
V15          -0.2422      0.092     -2.629      0.009     -0.423
-0.062
V2           -0.1791      0.062     -2.910      0.004     -0.300
-0.058
V16          -0.3186      0.065     -4.891      0.000     -0.446
-0.191
V3           -0.4383      0.086     -5.097      0.000     -0.607
-0.270
V8           -0.2271      0.045     -5.068      0.000     -0.315
```

-0.139					
V7	-0.1045	0.055	-1.888	0.059	-0.213
0.004					
V4	0.7933	0.088	8.965	0.000	0.620
0.967					
V13	-0.2437	0.093	-2.613	0.009	-0.427
-0.061					
V14	-0.8932	0.067	-13.258	0.000	-1.025
-0.761					
V22	0.2372	0.102	2.326	0.020	0.037
0.437					
V12	-0.7873	0.114	-6.891	0.000	-1.011
-0.563					
V21	0.1911	0.065	2.962	0.003	0.065
0.318					
V9	-0.2117	0.082	-2.592	0.010	-0.372
-0.052					

```
=====
=====
"""
```

```
logreg = LogisticRegression(C=1, penalty='l2')


logreg.fit(X_train[features], y_train)

LogisticRegression(C=1, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

predict_summarize(logreg.predict(X_train[features]), y_train, 0.8,
True)

Accuracy = 0.999332878553801
Sensitivity = 0.6942857142857143
Specificity = 0.9998693559247088
False Positive Rate = 0.00013064407529118555

Precision = 0.9033457249070632
Recall = 0.6942857142857143
Plotting
```

	predicted_no	predicted_yes
ind		
actual_no	198988	26
actual_yes	107	243

Observation:

This has a low recall value. We can continue with other models & recursive feature eliminations.

Random Forest

```

...
n_estimators = [50, 100]
max_features = ['auto', 'sqrt']
criterion = ["gini", "entropy"]
max_depth = [5, 10, 15]
min_samples_split = [30, 50]
min_impurity_decrease = [0.1, 0.2]
param_grid = {'n_estimators': n_estimators,
              'max_features': max_features,
              'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'criterion': criterion,

```

```

        'bootstrap':[True],
        'oob_score':[True]}

grid_search = GridSearchCV(estimator = RandomForestClassifier(),
                           param_grid = param_grid,
                           cv = kf, n_jobs = 8, verbose = 2)
grid_search.fit(X_train, y_train)
print("tuned hpyerparameters :(best parameters) ",
      grid_search.best_params_)
print("accuracy :", grid_search.best_score_)
'''

'\nn_estimators = [50, 100]\nmax_features = [\n'auto\n',\n'sqrt\n']\n
ncriterion = ["gini", "entropy"]\nmax_depth = [5,10,15]\n
nmin_samples_split = [30, 50]\nmin_impurity_decrease = [0.1, 0.2]\n
nparam_grid = {\n'n_estimators\n': n_estimators,\n
n                \n'max_features\n': max_features,\n
n                \n'max_depth\n': max_depth,\n
n                \n'min_samples_split\n': min_samples_split,\n
\n'criterion\n':criterion,\n
\n'bootstrap\n':[True],\n
\n'oob_score\n':[True]}\n\ngrid_search = GridSearchCV(estimator =
RandomForestClassifier(), \n
param_grid =
param_grid, \n
cv = kf, n_jobs = 8, verbose =
2)\ngrid_search.fit(X_train, y_train)\nprint("tuned hpyerparameters :
(best parameters) ", grid_search.best_params_)\nprint("accuracy :",
grid_search.best_score_)\n'

model_rf = RandomForestClassifier(bootstrap=True,
                                criterion = 'gini',
                                max_depth=5,
                                max_features='auto',
                                min_samples_split=50,
                                n_estimators=50,
                                min_impurity_decrease=0.1,
                                random_state = 42,
                                oob_score=True
                                )

model_rf.fit(X_train, y_train)


RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
class_weight=None,
                        criterion='gini', max_depth=5,
max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.1,
min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=50,
                        min_weight_fraction_leaf=0.0, n_estimators=50,
                        n_jobs=None, oob_score=True, random_state=42,
verbose=0,
                        warm_start=False)

```

```
predict_summarize([x[1] for x in model_rf.predict_proba(X_train)],  
y_train, 0.32, True)
```

```
Accuracy = 0.998244417246845  
Sensitivity = 0.0  
Specificity = 1.0  
False Positive Rate = 0.0
```

```
Precision = nan  
Recall = 0.0  
Plotting
```



	predicted_no	predicted_yes
ind		
actual_no	199014	0
actual_yes	350	0

Observation:

This model seems to be performing very poor by only predicting no.

RFE with Random Forest

```
model_rf = RandomForestClassifier(bootstrap=True,  
                                criterion = 'gini',
```

```
max_depth=5,  
max_features='auto',  
min_samples_split=50,  
n_estimators=50,  
min_impurity_decrease=0.1,  
random_state = 42,  
oob_score=True  
)
```


```
model_rf.fit(X_train[features], y_train)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,  
class_weight=None,  
criterion='gini', max_depth=5,  
max_features='auto',  
max_leaf_nodes=None, max_samples=None,  
min_impurity_decrease=0.1,  
min_impurity_split=None,  
min_samples_leaf=1, min_samples_split=50,  
min_weight_fraction_leaf=0.0, n_estimators=50,  
n_jobs=None, oob_score=True, random_state=42,  
verbose=0,  
warm_start=False)
```

```
predict_summarize([x[1] for x in  
model_rf.predict_proba(X_train[features])], y_train, 0.32, True)
```

```
Accuracy = 0.998244417246845  
Sensitivity = 0.0  
Specificity = 1.0  
False Positive Rate = 0.0
```

```
Precision = nan  
Recall = 0.0  
Plotting
```



	predicted_no	predicted_yes
ind		
actual_no	199014	0
actual_yes	350	0

Notice that the metrics is same even after using RFE features, hence not deleting features for final variable selection

Model building with balancing Classes

Perform class balancing with :

- Random Oversampling
- SMOTE
- ADASYN

This can be done on the train data to help the model recognise the fraudulent transactions better.

Class Imbalance

```
print('Fraud flag 1 count: ',y.value_counts()[1])
print('Fraud flag 0 count: ',y.value_counts()[0])
ClassImbRatio=y.value_counts()[1]/len(y) * 100
```

```
print('Class imbalance ratio: ',round(ClassImbRatio,3))
```

```
Fraud flag 1 count: 492  
Fraud flag 0 count: 284315  
Class imbalance ratio: 0.173
```

SMOTE

Implementing SMOTE

```
smote=SMOTE(sampling_strategy=0.3, random_state=42, k_neighbors=3)
```

```
X_train_smote,y_train_smote=smote.fit_sample(X_train,y_train)
```

```
print('Before SMOTE',Counter(y_train))  
print('After SMOTE',Counter(y_train_smote))
```

```
Before SMOTE Counter({0: 199014, 1: 350})  
After SMOTE Counter({0: 199014, 1: 59704})
```

ADASYN

```
X_train_adasyn, y_train_adasyn = ADASYN(sampling_strategy=0.3,  
random_state=42).fit_sample(X_train,y_train)
```

```
print('Before ADASYN',Counter(y_train))  
print('After ADASYN',Counter(y_train_adasyn))
```

```
Before ADASYN Counter({0: 199014, 1: 350})  
After ADASYN Counter({0: 199014, 1: 59666})
```

Print the class distribution after applying SMOTE

```
X_train_smote_1 = X_train_smote[X_train.shape[0]:]
```

```
X_train_1 = X_train.to_numpy()[numpy.where(y_train==1.0)]  
X_train_0 = X_train.to_numpy()[numpy.where(y_train==0.0)]
```



```
plt.rcParams['figure.figsize'] = [20, 20]  
fig = plt.figure()
```

```
plt.subplot(2, 1, 1)  
plt.scatter(X_train_1[:, 0], X_train_1[:, 1], label='Actual Class-1  
Examples')  
plt.scatter(X_train_smote_1.iloc[:X_train_1.shape[0], 0],  
X_train_smote_1.iloc[:X_train_1.shape[0], 1],  
label='Artificial SMOTE Class-1 Examples')  
plt.legend()
```

```
plt.subplot(2, 1, 2)  
plt.scatter(X_train_1[:, 0], X_train_1[:, 1], label='Actual Class-1  
Examples')  
plt.scatter(X_train_0[:X_train_1.shape[0], 0],
```

```
X_train_0[:X_train_1.shape[0], 1], label='Actual Class-0 Examples')
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7f5783f47890>
```



Print the class distribution after applying ADASYN

```
import warnings
```

```
warnings.filterwarnings("ignore")
```

```
# Artificial minority samples and corresponding minority labels from  
ADASYN are appended
```

```
# below X_train and y_train respectively
```

```
# So to exclusively get the artificial minority samples from ADASYN,  
we do
```

```
X_train_adasyn_1 = X_train_adasyn[X_train.shape[0]:]
```

```
X_train_1 = X_train.to_numpy()[numpy.where(y_train==1.0)]
```

```
X_train_0 = X_train.to_numpy()[numpy.where(y_train==0.0)]
```

```



import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = [20, 20]
fig = plt.figure()

plt.subplot(2, 1, 1)
plt.scatter(X_train_1[:, 0], X_train_1[:, 1], label='Actual Class-1
Examples')
plt.scatter(X_train_adasyn_1.iloc[:X_train_1.shape[0], 0],
X_train_adasyn_1.iloc[:X_train_1.shape[0], 1],
            label='Artificial ADASYN Class-1 Examples')
plt.legend()

plt.subplot(2, 1, 2)
plt.scatter(X_train_1[:, 0], X_train_1[:, 1], label='Actual Class-1
Examples')
plt.scatter(X_train_0[:X_train_1.shape[0], 0],
X_train_0[:X_train_1.shape[0], 1], label='Actual Class-0 Examples')
plt.legend()

<matplotlib.legend.Legend at 0x7f5783d24390>

```

Looking at the distribution of the classes, it seems like ADASYN has created more similar samples as compared to SMOTE

Models on Oversampled Data

Logistic Regression - SMOTE

```

'''
grid_params = {
    "C":numpy.logspace(-3, 3,7),
    "penalty":["l1", "l2"]
}
logreg = LogisticRegression()
logreg_cv = GridSearchCV(logreg, grid_params, cv=kf)
logreg_cv.fit(X_train_smote, y_train_smote)
print("tuned hyperparameters :(best parameters) ",
logreg_cv.best_params_)

```

```

print("accuracy :", logreg_cv.best_score_)
'''

'\ngrid_params = {\n    "C":numpy.logspace(-3, 3,7), \n    "penalty":
["l1", "l2"]\n}\n\nlogreg = LogisticRegression()\nlogreg_cv =
GridSearchCV(logreg, grid_params, cv=kf)\nlogreg_cv.fit(X_train_smote,
y_train_smote)\nprint("tuned hpyerparameters :(best parameters) ",
logreg_cv.best_params_)\nprint("accuracy :", logreg_cv.best_score_)\n'


logreg = LogisticRegression(C=10, penalty='l2')

logreg.fit(X_train_smote, y_train_smote)

LogisticRegression(C=10, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

plot_precision_recall_curve(logreg.predict(X_train_smote),
y_train_smote)

```



```

predict_summarize(logreg.predict(X_train_smote), y_train_smote, 0.9,
True)


```

```

Accuracy = 0.9714128897100318
Sensitivity = 0.8992697306713118
Specificity = 0.9930557649210608
False Positive Rate = 0.00694423507893917

```

Precision = 0.9749055781522371
Recall = 0.8992697306713118
Plotting



	predicted_no	predicted_yes
ind		
actual_no	197632	1382
actual_yes	6014	53690

Observation:

The recall value has improved upon balancing of data, we can try the same with adasyn & random forest.

Logistic Regression - ADASYN

```
...  
grid_params = {  
    "C":numpy.logspace(-3, 3,7),  
    "penalty":["l1", "l2"]  
}  
logreg = LogisticRegression()  
logreg_cv = GridSearchCV(logreg, grid_params, cv=kf)  
logreg_cv.fit(X_train_adasyn, y_train_adasyn)  
print("tuned hyperparameters :(best parameters) ",
```

```

logreg_cv.best_params_)
print("accuracy :", logreg_cv.best_score_)
'''

'\ngrid_params = {\n    "C":numpy.logspace(-3, 3,7), \n    "penalty":
["l1", "l2"]\n}\nlogreg = LogisticRegression()\nlogreg_cv =
GridSearchCV(logreg, grid_params, cv=kf)\n
logreg_cv.fit(X_train_adasyn, y_train_adasyn)\nprint("tuned
hyperparameters :(best parameters) ", logreg_cv.best_params_)\n
nprint("accuracy :", logreg_cv.best_score_)\n'


logreg = LogisticRegression(C=0.001, penalty='l2')

logreg.fit(X_train_adasyn, y_train_adasyn)

LogisticRegression(C=0.001, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

plot_precision_recall_curve(logreg.predict(X_train_adasyn),
y_train_adasyn)

```



```

predict_summarize(logreg.predict(X_train_adasyn), y_train_adasyn, 0.8,
True)

Accuracy = 0.9117867635688882
Sensitivity = 0.7005664867763886
Specificity = 0.9751123036570292


```

False Positive Rate = 0.024887696342970847

Precision = 0.8940602742070027

Recall = 0.7005664867763886

Plotting



	predicted_no	predicted_yes
ind		
actual_no	194061	4953
actual_yes	17866	41800

Observation:

The recall value was relatively better in case of smote + logistic.

RFE with SMOTE data for Logistics Regression

...

```
features=X_train_smote.columns
```

```
X_train_sm = sm.add_constant(X_train_smote[features])  
logm2 = sm.GLM(y_train_smote,X_train_sm, family =  
sm.families.Binomial())  
model = logm2.fit()
```

```

model.summary()
'''

'\nfeatures=X_train_smote.columns\n\nX_train_sm =
sm.add_constant(X_train_smote[features])\nlogm2 =
sm.GLM(y_train_smote,X_train_sm, family = sm.families.Binomial())\n
nmodel = logm2.fit()\nmodel.summary()\n'

# Removing features with p-value >0.05 & not removing Amount as it is
an Important feature

features=list(set(X_train_smote.columns)-set(['V23','V27']))


logreg = LogisticRegression(C=1, penalty='l2')

logreg.fit(X_train_smote[features], y_train_smote)

LogisticRegression(C=1, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, ll_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

plot_precision_recall_curve(logreg.predict(X_train_smote[features]),
y_train_smote)

```




```

predict_summarize(logreg.predict(X_train_smote[features]),
y_train_smote, 0.9, True)

```

Accuracy = 0.9714360809839284
Sensitivity = 0.8992697306713118
Specificity = 0.9930859135538204
False Positive Rate = 0.006914086446179666

Precision = 0.9750118040169978
Recall = 0.8992697306713118
Plotting



	predicted_no	predicted_yes
ind		
actual_no	197638	1376
actual_yes	6014	53690

Observation:

The recall value has not changed much as compared to the metric prior to rfe

RFE with ADASYN data

Removing features with p-value >0.05 & not removing Amount as it is an Important feature

```
features=list(set(X_train_adasyn.columns)-set(['V23','V27']))
```

```


logreg = LogisticRegression(C=1, penalty='l2')

logreg.fit(X_train_adasyn[features], y_train_adasyn)

LogisticRegression(C=1, class_weight=None, dual=False,
fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
verbose=0,
                    warm_start=False)

plot_precision_recall_curve(logreg.predict(X_train_adasyn[features]),
y_train_adasyn)

```



```

predict_summarize(logreg.predict(X_train_adasyn[features]),
y_train_adasyn, 0.9, True)

```

```


Accuracy = 0.9121037575382712
Sensitivity = 0.7149465357154828
Specificity = 0.9712130804867999
False Positive Rate = 0.028786919513200077

```

```

Precision = 0.8816004298675264
Recall = 0.7149465357154828
Plotting

```

	predicted_no	predicted_yes
ind		
actual_no	193285	5729
actual_yes	17008	42658

Observation:

The recall value has not changed much as compared to the metric prior to rfe

Random Forest - SMOTE

```
'''
n_estimators = [50, 100]
max_features = ['auto', 'sqrt']
criterion = ["gini", "entropy"]
max_depth = [5, 10, 15]
min_samples_split = [30, 50]
min_impurity_decrease = [0.1, 0.2]
param_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'criterion': criterion,
               'bootstrap': [True],
```

```

        'oob_score':[True]}

grid_search = GridSearchCV(estimator = RandomForestClassifier(),
                           param_grid = param_grid,
                           cv = kf, n_jobs = 8, verbose = 2)
grid_search.fit(X_train_smote, y_train_smote)
print("tuned hpyerparameters :(best parameters) ",
      grid_search.best_params_)
print("accuracy :", grid_search.best_score_)
'''

'\nn_estimators = [50, 100]\nmax_features = [\n'auto\n',\n'sqrt\n']\n
ncriterion = ["gini", "entropy"]\nmax_depth = [5,10,15]\n
nmin_samples_split = [30, 50]\nmin_impurity_decrease = [0.1, 0.2]\n
nparam_grid = {\n'n_estimators\n': n_estimators,\n
n               \n'max_features\n': max_features,\n
n               \n'max_depth\n': max_depth,\n
n               \n'min_samples_split\n': min_samples_split,\n
\n'criterion\n':criterion,\n
\n'bootstrap\n':[True],\n
\n'oob_score\n':[True]}\n\ngrid_search = GridSearchCV(estimator =
RandomForestClassifier(), \n
param_grid, \n
cv = kf, n_jobs = 8, verbose =
2)\ngrid_search.fit(X_train_smote, y_train_smote)\nprint("tuned
hpyerparameters :(best parameters) ", grid_search.best_params_)\n
print("accuracy :", grid_search.best_score_)\n'


model_rf_smote = RandomForestClassifier(bootstrap=True,
                                       criterion = 'entropy',
                                       max_depth=15,
                                       max_features='sqrt',
                                       min_samples_split=30,
                                       n_estimators=50,
                                       random_state = 42,
                                       oob_score=True
                                       )

model_rf_smote.fit(X_train_smote, y_train_smote)

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
class_weight=None,
                       criterion='entropy', max_depth=15,
max_features='sqrt',
                       max_leaf_nodes=None, max_samples=None,
                       min_impurity_decrease=0.0,
min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=30,
                       min_weight_fraction_leaf=0.0, n_estimators=50,
                       n_jobs=None, oob_score=True, random_state=42,
verbose=0,
                       warm_start=False)

```


```
plot_precision_recall_curve([x[1] for x in  
model_rf_smote.predict_proba(X_train_smote)], y_train_smote.values)
```



```
predict_summarize([x[1] for x in  
model_rf_smote.predict_proba(X_train_smote)], y_train_smote, 0.5,  
True)
```

Accuracy = 0.9996444004669176
Sensitivity = 0.9988610478359908
Specificity = 0.999879405468962
False Positive Rate = 0.00012059453103801742

Precision = 0.999597720415689
Recall = 0.9988610478359908
Plotting



	predicted_no	predicted_yes
ind		
actual_no	198990	24
actual_yes	68	59636

Observation:

There is a drastic improvement in the recall value as per the balanced train data is concerned, we can try another iteration with adasyn sampled data.

Random Forest - ADASYN

```

...
n_estimators = [50, 100]
max_features = ['auto', 'sqrt']
criterion = ["gini", "entropy"]
max_depth = [5, 10, 15]
min_samples_split = [30, 50]
min_impurity_decrease = [0.1, 0.2]
param_grid = {'n_estimators': n_estimators,
              'max_features': max_features,
              'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'criterion': criterion,

```

```

        'bootstrap':[True],
        'oob_score':[True]}

grid_search = GridSearchCV(estimator = RandomForestClassifier(),
                           param_grid = param_grid,
                           cv = kf, n_jobs = 8, verbose = 2)
grid_search.fit(X_train_adasyn, y_train_adasyn)
print("tuned hpyerparameters :(best parameters) ",
      grid_search.best_params_)
print("accuracy :", grid_search.best_score_)
'''

'\nn_estimators = [50, 100]\nmax_features = [\n'auto\n',\n'sqrt\n']\n
ncriterion = ["gini", "entropy"]\nmax_depth = [5,10,15]\n
nmin_samples_split = [30, 50]\nmin_impurity_decrease = [0.1, 0.2]\n
nparam_grid = {\n'n_estimators\n': n_estimators,\n
n              \n'max_features\n': max_features,\n
n              \n'max_depth\n': max_depth,\n
n              \n'min_samples_split\n': min_samples_split,\n
\n'\n'criterion\n':criterion,\n
\n              \n'\n'bootstrap\n':[True],\n
\n              \n'\n'oob_score\n':[True]}\n\ngrid_search = GridSearchCV(estimator =
RandomForestClassifier(), \n
                           param_grid =
param_grid, \n
                           cv = kf, n_jobs = 8, verbose =
2)\ngrid_search.fit(X_train_adasyn, y_train_adasyn)\nprint("tuned
hpyerparameters :(best parameters) ", grid_search.best_params_)\n
nprint("accuracy :", grid_search.best_score_)\n'


model_rf_adasyn = RandomForestClassifier(bootstrap=True,
                                         criterion = 'entropy',
                                         max_depth=15,
                                         max_features='sqrt',
                                         min_samples_split=30,
                                         n_estimators=50,
                                         random_state = 42,
                                         oob_score=True
                                         )

model_rf_adasyn.fit(X_train_adasyn, y_train_adasyn)

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
class_weight=None,
                           criterion='entropy', max_depth=15,
max_features='sqrt',
                           max_leaf_nodes=None, max_samples=None,
                           min_impurity_decrease=0.0,
min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=30,
                           min_weight_fraction_leaf=0.0, n_estimators=50,
                           n_jobs=None, oob_score=True, random_state=42,
verbose=0,
                           warm_start=False)

```


```
plot_precision_recall_curve([x[1] for x in
model_rf_adasyn.predict_proba(X_train_adasyn)], y_train_adasyn.values)
```



```
predict_summarize([x[1] for x in
model_rf_adasyn.predict_proba(X_train_adasyn)], y_train_adasyn, 0.6,
True)
```

```
Accuracy = 0.999845368795423
Sensitivity = 0.9999162001810076
Specificity = 0.9998241329755696
False Positive Rate = 0.0001758670244304421
```

```
Precision = 0.9994136960600375
Recall = 0.9999162001810076
Plotting
```



	predicted_no	predicted_yes
ind		
actual_no	198979	35
actual_yes	5	59661


Observation:

The recall value has further improved and by far is the best in case of adasyn + random forest. We can retain this to be the final model & check the recall as per the test data.

```
predict_summarize([x[1] for x in
model_rf_adasyn.predict_proba(X_test)], y_test, 0.6, True)
```

```
Accuracy = 0.9992392589211521
Sensitivity = 0.795774647887324
Specificity = 0.9995779650883343
False Positive Rate = 0.0004220349116657483
```

```
Precision = 0.7583892617449665
Recall = 0.795774647887324
Plotting
```




	predicted_no	predicted_yes
ind		
actual_no	85265	36
actual_yes	29	113

Observation:

The recall value on the test data is quite different from the train data, but for a data with imbalance as low as 0.17%, the recall value of 80% is quite good.

```
plot_feature_importance(X_train_adasyn.columns,
model_rf_adasyn.feature_importances_)
```

The top influencing parameters are

- V14
- V4
- V17
- V10
- V12

This can help reduce the fraudulent transaction significantly as it reduces the number of transactions to be underwritten or verified, thus reducing the effort for the same process while increasing the safety. Only the transactions which are flagged by the model can be verified.